

Course: DD2424 - Assignment 1

Exercise 1: Training a multi-linear classifier

Álvaro Orgaz Expósito

April 14, 2019

For this assignment I will just use data in the CIFAR 10 file *data batch 1* for training, the file *data batch 2* for validation and the file *test batch* for testing. Each batch contains 10000 images (columns) of 3072 pixels (rows) and 10 possible labels (targets). In figure 1 we can see the first 5 images of each label for *batch test*. And in figure 2 we can find the distribution of labels by images batch of CIFAR 10. Importantly the class distribution in training validation and test is approximately uniform.

I will convert the images data to double format and divide it by 255 to convert the data to be between 0 and 1 since normalization helps when training the networks.

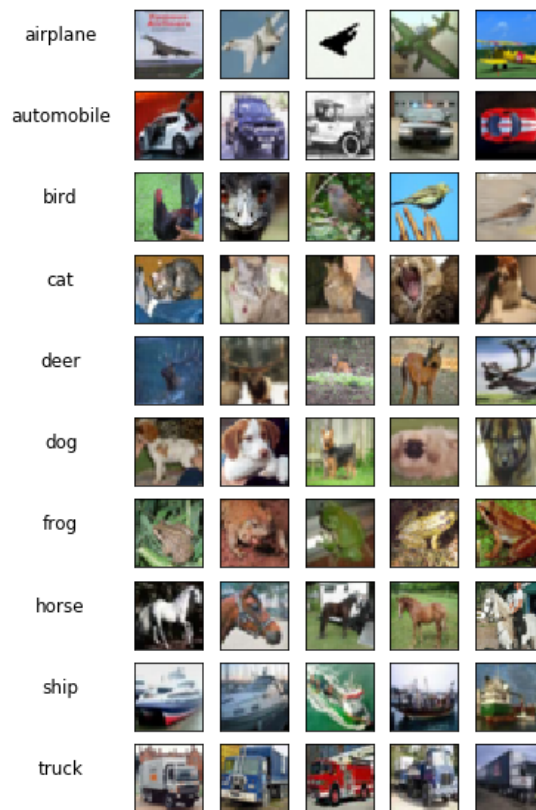


Figure 1: First 5 images of each label in the batch test.

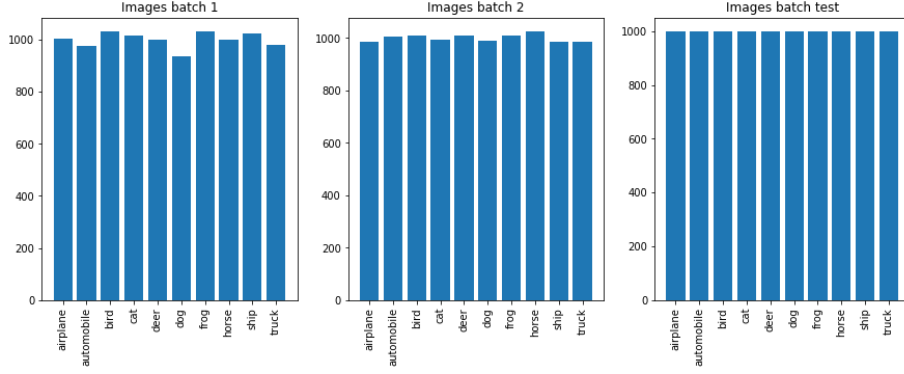


Figure 2: Distribution of labels by images batch of CIFAR 10.

In this assignment, I will use a neural network with one input layer (as many nodes as images pixels) and one output layer (as many nodes as possible image labels) with the loss function cross entropy (classification metric) and an L2 regularization term on the weight matrix. It is a simple network where no hidden layer will be used and the predicted class corresponds to the label with the highest predicted probability. Then, the shape of the weights matrix is as many rows as output nodes and as many columns as input nodes. And for the bias only 1 column with as many rows as output nodes.

Then, for computing the gradients of the cross entropy loss function and an L2 regularization term in mini-batch mode of gradient descent (back propagation) I have used the following equations by data batches, where D is the number of input features, N is the number of training samples and C is the number of possible output labels.

- Initial input and output data:

$$\begin{matrix} X & Y \\ [D \times N] & [C \times N] \end{matrix} \quad (1)$$

- Forward pass (predictions):

$$\begin{matrix} S \\ [C \times N] \end{matrix} = \begin{matrix} W \\ [C \times D] \end{matrix} \times \begin{matrix} X \\ [D \times N] \end{matrix} + \begin{matrix} b \\ [C \times 1] \end{matrix} \quad (2)$$

$$\begin{matrix} P \\ [C \times N] \end{matrix} = SoftMax \left(\begin{matrix} S \\ [C \times N] \end{matrix} \right) \quad where \quad SoftMax(S_{ji}) = \frac{exp(S_{ji})}{\sum_{c=1}^C exp(S_{ci})} \quad (3)$$

- Loss function (cross entropy) to minimise respect to weights:

$$J = L + \lambda ||W||^2 = \frac{1}{N} \sum_{i=1}^N -\log \left(\begin{matrix} Y_{:i}^T \\ [1 \times C] \end{matrix} \times \begin{matrix} P_{:i} \\ [C \times 1] \end{matrix} \right) + \lambda \sum_{c,d}^{C,D} W_{cd}^2 \quad (4)$$

- Backward pass (vanilla updates equal to $-\eta$ times the following gradients):

$$\frac{\partial \lambda ||W||^2}{\partial W} = 2\lambda \begin{matrix} W \\ [C \times D] \end{matrix} \quad and \quad \frac{\partial L}{\partial W} = \frac{1}{N} \left(\begin{matrix} P \\ [C \times N] \end{matrix} - \begin{matrix} Y \\ [C \times N] \end{matrix} \right) \times \begin{matrix} X^T \\ [N \times D] \end{matrix} \quad (5)$$

$$then \quad \frac{\partial J}{\partial W} = \frac{\partial L}{\partial W} + \frac{\partial \lambda ||W||^2}{\partial W} \quad and \quad \frac{\partial J}{\partial b} = \frac{\partial L}{\partial b} = \frac{1}{N} \left(\begin{matrix} P \\ [C \times N] \end{matrix} - \begin{matrix} Y \\ [C \times N] \end{matrix} \right) \times \begin{matrix} 1 \\ [N \times 1] \end{matrix} \quad (6)$$

I want to check that the gradients computed in the gradients function correspond to the correct gradients. To do this, I compare the gradient obtained by the network (analytical method) with the gradient computed with the finite difference method (numerical method), after initializing the weight and bias with same Gaussian random (zero mean and standard deviation 0.01) for both cases. As error measures I will apply the absolute error and relative error in two different scenarios:

- Firstly, I study the gradients of the 20 first dimensions of the first image of *data batch 1* without regularization. The result is:
 - For weights W: 100.0% of absolute errors below 1e-6 and the maximum absolute error is 1.5539e-08
 - For bias b: 100.0% of absolute errors below 1e-6 and the maximum absolute error is 4.6837e-08
 - For weights W: 100.0% of relative errors below 1e-6 and the maximum relative error is 1.3081e-07
 - For bias b: 100.0% of relative errors below 1e-6 and the maximum relative error is 2.2443e-07
- Secondly, I study the gradients of all dimensions of the first 5 images of *data batch 1* with regularization. The result is:
 - For weights W: 100.0% of absolute errors below 1e-6 and the maximum absolute error is 4.7962e-08
 - For bias b: 100.0% of absolute errors below 1e-6 and the maximum absolute error is 5.8252e-08
 - For weights W: 92.8255% of relative errors below 1e-6 and the maximum relative error is 0.0096
 - For bias b: 100.0% of relative errors below 1e-6 and the maximum relative error is 4.4302e-07

To evaluate the learning process I use the metrics loss and accuracy (fraction of well-predicted images or images with the highest predicted class equal to the true) after each epoch of the vanilla mini-batch gradient descent algorithm. In figure 3 we can find the learning curves (both metrics for training and validation sets) for the 4 parametrization cases proposed.

$$W^{t+1} = W^t - \eta \frac{\partial J(\text{Batch}^{t+1})}{\partial W} \quad \text{where} \quad \frac{\partial J(\text{Batch}^{t+1})}{\partial W} = \frac{1}{N_{\text{Batch}^{t+1}}} \sum_{i=1}^{N_{\text{Batch}^{t+1}}} \frac{\partial J_i}{\partial W} \quad (7)$$

$$b^{t+1} = b^t - \eta \frac{\partial J(\text{Batch}^{t+1})}{\partial b} \quad \text{where} \quad \frac{\partial J(\text{Batch}^{t+1})}{\partial b} = \frac{1}{N_{\text{Batch}^{t+1}}} \sum_{i=1}^{N_{\text{Batch}^{t+1}}} \frac{\partial J_i}{\partial b} \quad (8)$$

On the one side, a large learning rate ($\eta=0.1$) leads to instability in the training process and the loss does not seem to converge. On the other side, regularization increases training error and simultaneously reduces the gap between training and validation error curves. In this case, the net effect is always a higher validation error if I increase λ which is obviously not what I want but definitely what we can expect given that ridge regression usually only makes sense if we train models with high capacity on relatively small datasets.

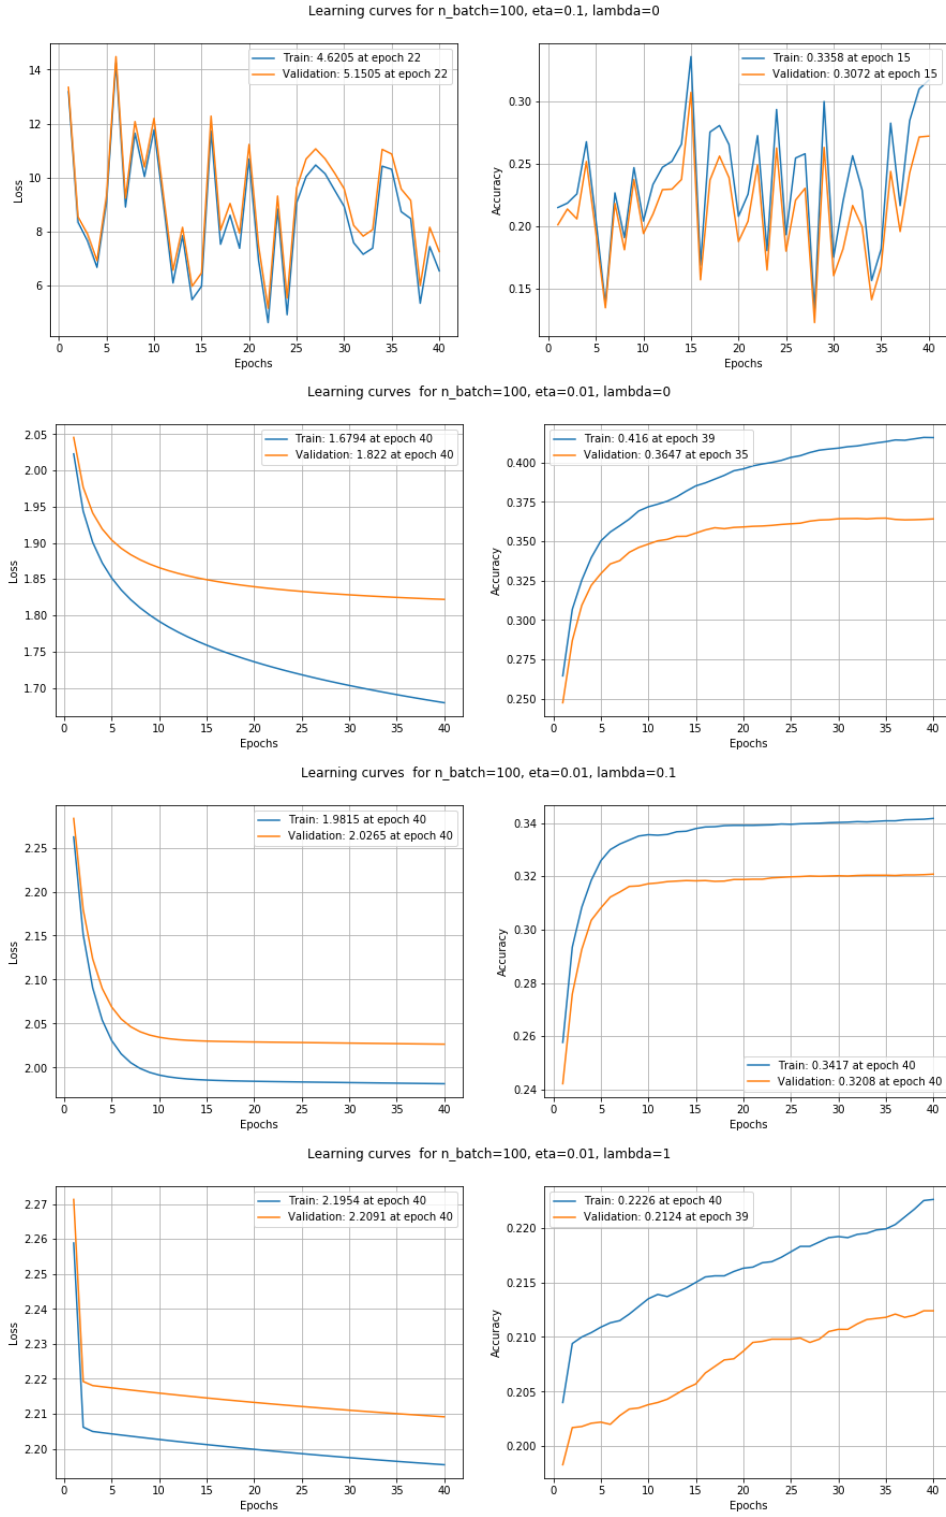


Figure 3: Learning curves for each 4 parametrization case.

In figure 4 we can see what class templates the network has learnt for each output node (weights matrix by rows) and parametrization case. Notice how regularization smooths out the weights, this is to be expected since ridge regression forces the weights to be closer to each other in value (i.e. closer to zero). Also, for some of the classes, we can see the shape of the image or contour of objects described such as for automobile or horse. Also, in figure 5 we can observe the histograms of weights by regularizations where the largest regularizations λ have a lower standard deviation (more centred to 0).




Figure 4: Weights learnt for each 4 parametrization case.




Figure 5: Histogram of weights by regularization levels (cases 2, 3 and 4).

In figure 6 we can find the test accuracy and confusion matrix for each parametrization case. Interestingly the networks with low overall accuracy seem to assign mostly a single class to most test samples. Notice that even the best network has a tendency to classify all animals like cats.

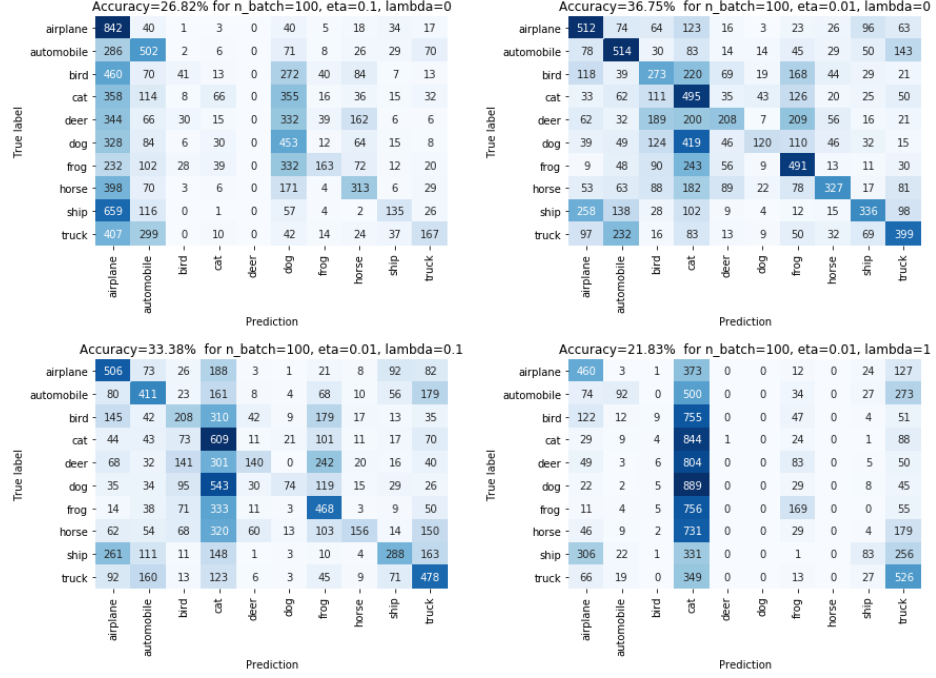


Figure 6: Test accuracy and confusion matrix for each parametrization case.