

## Course: DD2424 - Assignment 3

Álvaro Orgaz Expósito

April 29, 2019

In this assignment, we will use a  $k$ -layer neural network with one input layer (as many nodes as images pixels),  $k-1$  hidden layer (with *ReLU* or rectified linear units as activation function), and one output layer (as many nodes as possible image labels and *SoftMax* as activation function) with the loss function cross entropy (classification metric) with L2 regularization. Then, the predicted class corresponds to the label with the highest predicted probability (since the sum of the output activations for a sample is 1 as a consequence of using *SoftMax*). The shape of the weights matrix is as many rows as output nodes and as many columns as input nodes for each layer. And for the bias, gammas and betas only 1 column with as many rows as layer nodes (the last 2 parameters for the *Batch Normalization* implemented).

Then, for computing a forward pass through the network or prediction, given some parameters and *Batch Normalization* in hidden layers, we use from the initial input image  $X^0$ :

for  $l = 1, \dots, k - 1$

$$S^l = W^l \times X^{l-1} + b^l \quad (1)$$

$$\hat{S}^l = \text{BatchNormalize}(S^l, \mu^l, \sigma^l) \quad (2)$$

$$\tilde{S}^l = \gamma^l \odot \hat{S}^l + \beta^l \quad (3)$$

$$X^l = \text{ReLU}(\tilde{S}^l) \quad \text{where} \quad \text{ReLU}(\hat{S}_{ij}^l) = \max(0, \hat{S}_{ij}^l) \quad (4)$$

and then finally

$$S = W^k \times X^{k-1} + b^k \quad (5)$$

$$P = \text{SoftMax}(S) \quad \text{where} \quad \text{SoftMax}(S_{ji}) = \frac{\exp(S_{ji})}{\sum_{c=1}^C \exp(S_{ci})} \quad (6)$$

where

$$\text{BatchNormalize}(S^l, \mu^l, \sigma^l) = \text{diag}(\sigma^l + \epsilon)^{-1} (S^l - \mu^l) \quad (7)$$

and  $\epsilon$  is a small number, and  $S^l$  and  $\tilde{S}^l$ , which are the un/normalized scores, have as many rows as layer nodes (as well as  $\mu^l$  and  $\sigma^l$ ).

For computing the mean and standard deviation of each dimension or layer node in the batch of size  $n$  we use

$$\mu^l = \frac{\sum_{i=1}^n S_i^l}{n} \quad (8)$$

$$\sigma_j^l = \sqrt{\frac{\sum_{i=1}^n (S_{ji}^l - \mu_j^l)^2}{n}} \quad \text{for } j = 1, \dots, \# \text{ layer nodes} \quad (9)$$

but note that for validating and testing we use the means and standard deviations calculated during the training.

Also, the loss function (cross entropy) to minimise respect to parameters is:

$$J = L + \lambda \sum_{l=1}^k \|W^l\|^2 = \frac{1}{N} \sum_{i=1}^N -\log \left( \begin{matrix} Y_{:i}^T \\ [1 \times C] \end{matrix} \times \begin{matrix} P_{:i} \\ [C \times 1] \end{matrix} \right) + \lambda \sum_{l=1}^k \|W^l\|^2 \quad (10)$$

Then, for computing the gradients of the cross-entropy loss function and an L2 regularization term in mini-batch mode (given some images and labels, the weights and the regularization parameter) I use the equations mentioned in the assignment instructions by data batches and then update the parameters of each layer ( $W^l$ ,  $b^l$ ,  $\gamma^l$  and  $\beta^l$ ) with updates equal to  $-\eta_{update}$  times the respective gradients.

$$W_{t+1}^k = W_t^k - \eta_{t+1} \frac{\partial J(Batch^{t+1})}{\partial W^k} \quad \text{where} \quad \frac{\partial J(Batch^{t+1})}{\partial W^k} = \frac{1}{N_{Batch^{t+1}}} \sum_{i=1}^{N_{Batch^{t+1}}} \frac{\partial J_i}{\partial W^k} \quad (11)$$

$$b_{t+1}^k = b_t^k - \eta_{t+1} \frac{\partial J(Batch^{t+1})}{\partial b^k} \quad \text{where} \quad \frac{\partial J(Batch^{t+1})}{\partial b^k} = \frac{1}{N_{Batch^{t+1}}} \sum_{i=1}^{N_{Batch^{t+1}}} \frac{\partial J_i}{\partial b^k} \quad (12)$$

$$\gamma_{t+1}^k = \gamma_t^k - \eta_{t+1} \frac{\partial J(Batch^{t+1})}{\partial \gamma^k} \quad \text{where} \quad \frac{\partial J(Batch^{t+1})}{\partial \gamma^k} = \frac{1}{N_{Batch^{t+1}}} \sum_{i=1}^{N_{Batch^{t+1}}} \frac{\partial J_i}{\partial \gamma^k} \quad (13)$$

$$\beta_{t+1}^k = \beta_t^k - \eta_{t+1} \frac{\partial J(Batch^{t+1})}{\partial \beta^k} \quad \text{where} \quad \frac{\partial J(Batch^{t+1})}{\partial \beta^k} = \frac{1}{N_{Batch^{t+1}}} \sum_{i=1}^{N_{Batch^{t+1}}} \frac{\partial J_i}{\partial \beta^k} \quad (14)$$

## 1 Data used & initialization of the parameters of the network

For this assignment I will just use data in the CIFAR 10, concretely 45.000 images from all 5 batch files *data batch #* for training, the remaining 5.000 images for validation and the file *test batch* for testing. Each batch contains 10000 images (columns) of 3072 pixels (rows) and 10 possible labels (targets). In figure 1 we can find the distribution of labels by created sets of CIFAR 10. Importantly the distribution in training validation and test is approximately uniform.

I will normalize (minus the mean and divided by standard deviation) the images by pixel (which are between 0 and 255) using the mean and standard deviation of the training data, since normalization helps when training the networks.

I will initialize the network parameters with the weights matrix with a Gaussian random distribution (zero mean and standard deviation  $2/\sqrt{\text{input dimension}}$  following the *He initialization*, the bias and  $\beta$  matrix with zeros, and  $\gamma$  with ones.




Figure 1: Distribution of labels by images sets of CIFAR 10.

## 2 Checking the computed gradients for the network parameters (WITHOUT *Batch Normalization*)

I want to check that the gradients computed in the gradients function without *Batch Normalization* correspond to the correct gradients. To do this, I compare the gradient obtained by the network (analytical method) with the gradient computed with the finite difference method (numerical method), after initializing the parameters as mentioned. As error measures I will apply the absolute error in 3 different scenarios:

- Firstly, for 2-layer network with 50 hidden nodes, I study the gradients of the 10 first dimensions of the first 5 images of the training set without regularization. The result is:
  - For weights  $W^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $2.9144e-08$ .
  - For bias  $b^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $1.9984e-08$ .
  - For weights  $W^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $1.0063e-06$ .
  - For bias  $b^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $8.0699e-08$ .
- Secondly, for 3-layer network with 50/50 hidden nodes, I study the gradients of the 10 first dimensions of the first 5 images of the training set without regularization. The result is:
  - For weights  $W^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $2.0314e-08$ .
  - For bias  $b^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $1.7173e-08$ .
  - For weights  $W^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $3.6173e-07$ .
  - For bias  $b^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $2.5908e-08$ .
  - For weights  $W^3$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $6.4744e-07$ .
  - For bias  $b^3$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $6.0997e-08$ .

- Thirdly, for 4-layer network with 50/50/25 hidden nodes, I study the gradients of the 10 first dimensions of the first 5 images of the training set without regularization. The result is:

- For weights  $W^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $2.237e-08$ .
- For bias  $b^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $1.7213e-08$
- For weights  $W^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $2.4880e-07$ .
- For bias  $b^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $1.8120e-08$ .
- For weights  $W^3$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $2.6750e-07$ .
- For bias  $b^3$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $4.6682e-08$ .
- For weights  $W^4$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $1.8688e-07$ .
- For bias  $b^4$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $7.9988e-08$ .

### 3 Checking the computed gradients for the network parameters (WITH *Batch Normalization*)

I want to check that the gradients computed in the gradients function with *Batch Normalization* correspond to the correct gradients. To do this, I compare the gradient obtained by the network (analytical method) with the gradient computed with the finite difference method (numerical method), after initializing the parameters as mentioned. As error measures I will apply the absolute error in 2 different scenarios:

- Firstly, for 2-layer network with 50 hidden nodes, I study the gradients of the 10 first dimensions of the first 5 images of the training set without regularization. The result is:
  - For weights  $W^1$ : 97.8% of absolute errors below  $1e-6$  and the maximum is  $1.7628e-06$ .
  - For bias  $b^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $1.3322e-16$ .
  - For weights  $\gamma^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $1.4987e-08$ .
  - For bias  $\beta^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $2.0604e-08$ .
  - For weights  $W^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $6.9411e-08$ .
  - For bias  $b^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $6.6373e-08$ .
- Secondly, for 3-layer network with 50/50 hidden nodes, I study the gradients of the 10 first dimensions of the first 5 images of the training set without regularization. The result is:
  - For weights  $W^1$ : 99.6% of absolute errors below  $1e-6$  and the maximum is  $1.3681e-06$ .
  - For bias  $b^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $8.8817e-17$ .
  - For weights  $\gamma^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $5.179e-08$ .
  - For bias  $\beta^1$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $6.1067e-08$ .
  - For weights  $W^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $3.9499e-07$ .
  - For bias  $b^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $4.9960e-17$ .
  - For weights  $\gamma^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $2.7055e-08$ .
  - For bias  $\beta^2$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $1.3721e-08$ .
  - For weights  $W^3$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $8.375e-08$ ].
  - For bias  $b^3$ : 100.0% of absolute errors below  $1e-6$  and the maximum is  $6.9580e-08$ .

## 4 Train k-layer network with cyclical learning rates

There is not really one optimal learning rate when training a neural network with vanilla (fixed  $\eta$ ) mini-batch gradient descent. Choose a too small learning rate and training will take too long and too large a learning rate may result in training diverging. For this assignment, I will explore the rather recent idea of exploiting cyclical learning rates *Smith, 2015* [1] as this approach eliminates much of the trial-and-error associated with finding a good learning rate and some of the costly hyperparameter optimization over multiple parameters associated with training with momentum. The main idea of cyclical learning rates is that during training the learning rate is periodically changed in a systematic fashion from a small value to a large one and then from this large value back to the small value. And this process is then repeated again and again until training is stopped. This assignment uses triangular CLR and the training is run for a set number of complete cycles and is stopped when the learning rate is at its smallest.

To evaluate the learning process I use the metrics loss and accuracy (fraction of well-predicted images or images with the highest predicted class equal to the true) after each epoch of the mini-batch gradient descent algorithm with cyclical learning rate (CLR). Note that I do not save the metrics after each update since evaluating the metrics for each batch makes the plots not smooth.

### 4.1 3-layer network (50 hidden nodes)

In figures 2 and 3 we can find the learning curves for the default parametrization case with and without *Batch Normalization*: batch size 100, 2 cycles,  $\eta_{min}$  1e-5,  $\eta_{max}$  1e-1, step size  $5 \cdot 45000 / 100 = 2250$ ,  $\lambda$  0.005. Note that it is equivalent to 20 epochs equal to 2 cycles  $\cdot$  2  $\cdot$  step size 2250 / (45000 images / batch size 100).

Also, in figures 4 and 5 we can find its test accuracy and confusion matrix which achieves an accuracy of 53.11% in the case without and 53.64% in the case with *Batch Normalization*.

### 4.2 9-layer network (50, 30, 20, 20, 10, 10, 10, and 10 hidden nodes)

In figures 6 and 7 we can find the learning curves for the default parametrization case with and without *Batch Normalization*: batch size 100, 2 cycles,  $\eta_{min}$  1e-5,  $\eta_{max}$  1e-1, step size  $5 \cdot 45000 / 100 = 2250$ ,  $\lambda$  0.005. Note that it is equivalent to 20 epochs equal to 2 cycles  $\cdot$  2  $\cdot$  step size 2250 / (45000 images / batch size 100).

Also, in figures 8 and 9 we can find its test accuracy and confusion matrix which achieves an accuracy of 50.32% in the case without and 52.97% in the case with *Batch Normalization*.

The conclusion after these results is that, at least with the default parametrization, increasing the number of hidden layer to 9 does not achieve better results. Also, the *Batch Normalization* technique helps the deeper network to improve the results, and this difference is not noticed in the network with only 3 hidden layers.

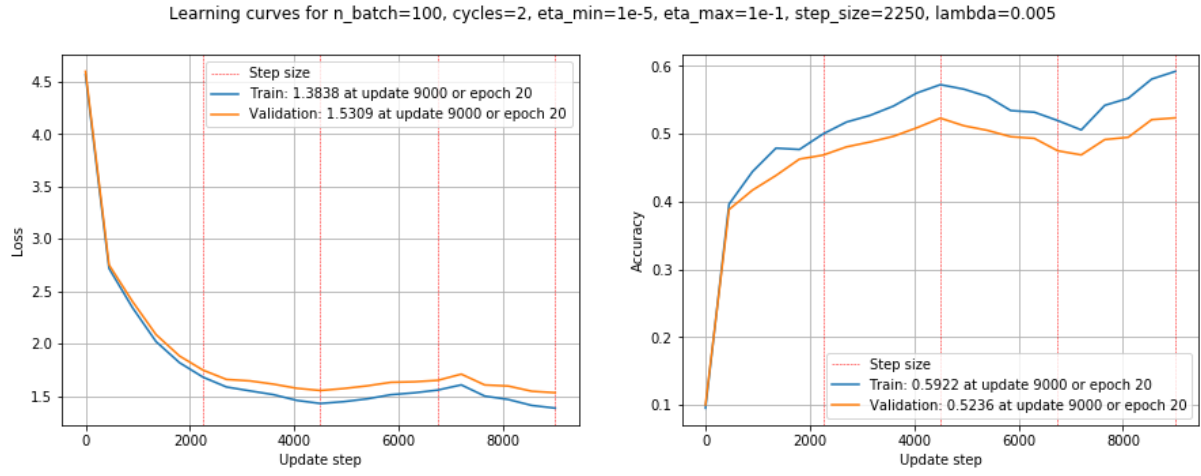


Figure 2: Learning curves for the default parametrization without *Batch Normalization* and 3 layers.




Figure 3: Learning curves for the default parametrization with *Batch Normalization* and 3 layers.




Figure 4: Test accuracy and confusion matrix for the default parametrization without *Batch Normalization* and 3 layers.




Figure 5: Test accuracy and confusion matrix for the default parametrization with *Batch Normalization* and 3 layers.




Figure 6: Learning curves for the default parametrization without *Batch Normalization* and 9 layers.




Figure 7: Learning curves for the default parametrization with *Batch Normalization* and 9 layers.




Figure 8: Test accuracy and confusion matrix for the default parametrization without *Batch Normalization* and 9 layers.




Figure 9: Test accuracy and confusion matrix for the default parametrization with *Batch Normalization* and 9 layers.

## 5 Coarse random search to set lambda

Firstly, I perform a coarse search over a very broad range of values for lambda to optimize the performance of the 3-layer network trained with *Batch Normalization*. About the range of  $\lambda$  values, I have tried a list of 10 values coming from  $10^{\text{random value}}$  with uniform sampling  $\text{random value} \sim \text{Uniform}(-5, -1)$ . This is the ordered list of  $\lambda$  for the coarse search: [0.00034, 0.00049, 0.00056, 0.00151, 0.00156, 0.00257, 0.00383, 0.00725, 0.03690, 0.07155].

In figure 10 we can find the output of the coarse search for  $\lambda$  regularization parameter. Here there is the performance of each  $\lambda$  value for the default parametrization case: batch size 100, 2 cycles,  $\eta_{\min}$  1e-5,  $\eta_{\max}$  1e-1, step size 5 epochs or 5 times number of images divided by batch size (in our case  $5 \cdot 45000 / 100 = 2250$ ). The conclusion is that within the range of values searched, we find a region of optimal values for the validation accuracy but not for the loss since the lower regularization value the better for this range searched (though not a big difference). Thus, choosing the optimal  $\lambda = 0.0015$ , we can see in figure 11 that the test accuracy achieved by this optimized 3-layer network with *Batch Normalization* is 53% which is not an improvement with respect to the previous test accuracy with default  $\lambda = 0.005$ .

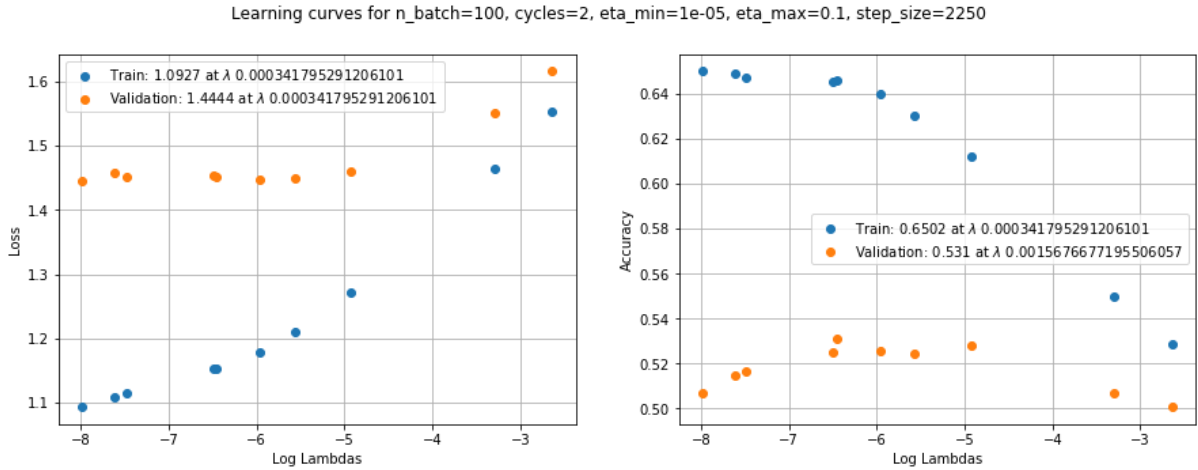


Figure 10: Coarse search for  $\lambda$  regularization parameter.




Figure 11: Test accuracy and confusion matrix for the default parametrization with *Batch Normalization* and optimal regularization.



## 6 Sensitivity to initialization

The previous results seem to indicate that the batch normalization is a good thing! The frequently stated pros of using it are that training becomes more stable, higher learning rates can be used and it acts as a form of regularization. Then, let's try different training regimes for initializing the parameters, instead of using *He initialization*, by using normally distributed initialization with 0 mean and  $\sigma$  equal to the same value at each layer. For three runs I set  $\sigma$  as 1e-1, 1e-3 and 1e-4 respectively and train the network with and without *Batch Normalization*.

In table 1 we can find the test accuracy for each experiment. It is very interesting how using *Batch Normalization* helps a lot to increase the performance or training with stability the networks with lower regularization term.

Sigma	Batch Normalization	Test Accuracy
0.1	0.0	0.5296
0.1	1.0	0.5378
0.001	0.0	0.1
0.001	1.0	0.536
0.0001	0.0	0.1
0.0001	1.0	0.536

Table 1: Table of test accuracies by sensitivity to initialization experiments.

## References

- [1] Smith, L. N. (2015). Cyclical learning rates for training neural networks. *arXiv:1506.01186*.